Methods 3: Multilevel Statistical Modeling and Machine Learning

Week 3: Generalized linear mixed effects models

September 28, 2021

by: Lau Møller Andersen

These slides are distributed according to the CC BY 4.0 licence:

https://creativecommons.org/licenses/by/4.0/

Attribution 4.0 International (CC BY 4.0)

Messages

- Practical exercise due 23.59 tomorrow

- Make sure to add your GitHub repository – a few are still missing:

  https://cryptpad.fr/pad/#/2/pad/edit/U21qNTbLgf kRiGZU1bnmDE2o/

- Remember, Class 2 (10-12) will be in 1453-116 tomorrow

RECAP *on pooling*

SLEEP STUDY EXAMPLE

https://psyteachr.github.io/stat-models-v1/introducing-linear-mixed-effects-models.html

Learning goals and outline –

Linear Mixed Effects Models (LMM)

1) Why can it be a good idea to do mixed effects modelling?

2) Understanding the basics of multilevel modelling

  - also known as linear mixed effects modelling

3) Appreciating the difference between the different levels of effects

  – or *random* and *fixed* effects, as they are also called

4) Understanding the concept of pooling (none, complete and partial)

Pooling - Summary

- Complete pooling
  - ignoring a categorical predictor (e.g. *subject*)

- No pooling
  - model each level of the categorical predictor separately

- Partial pooling
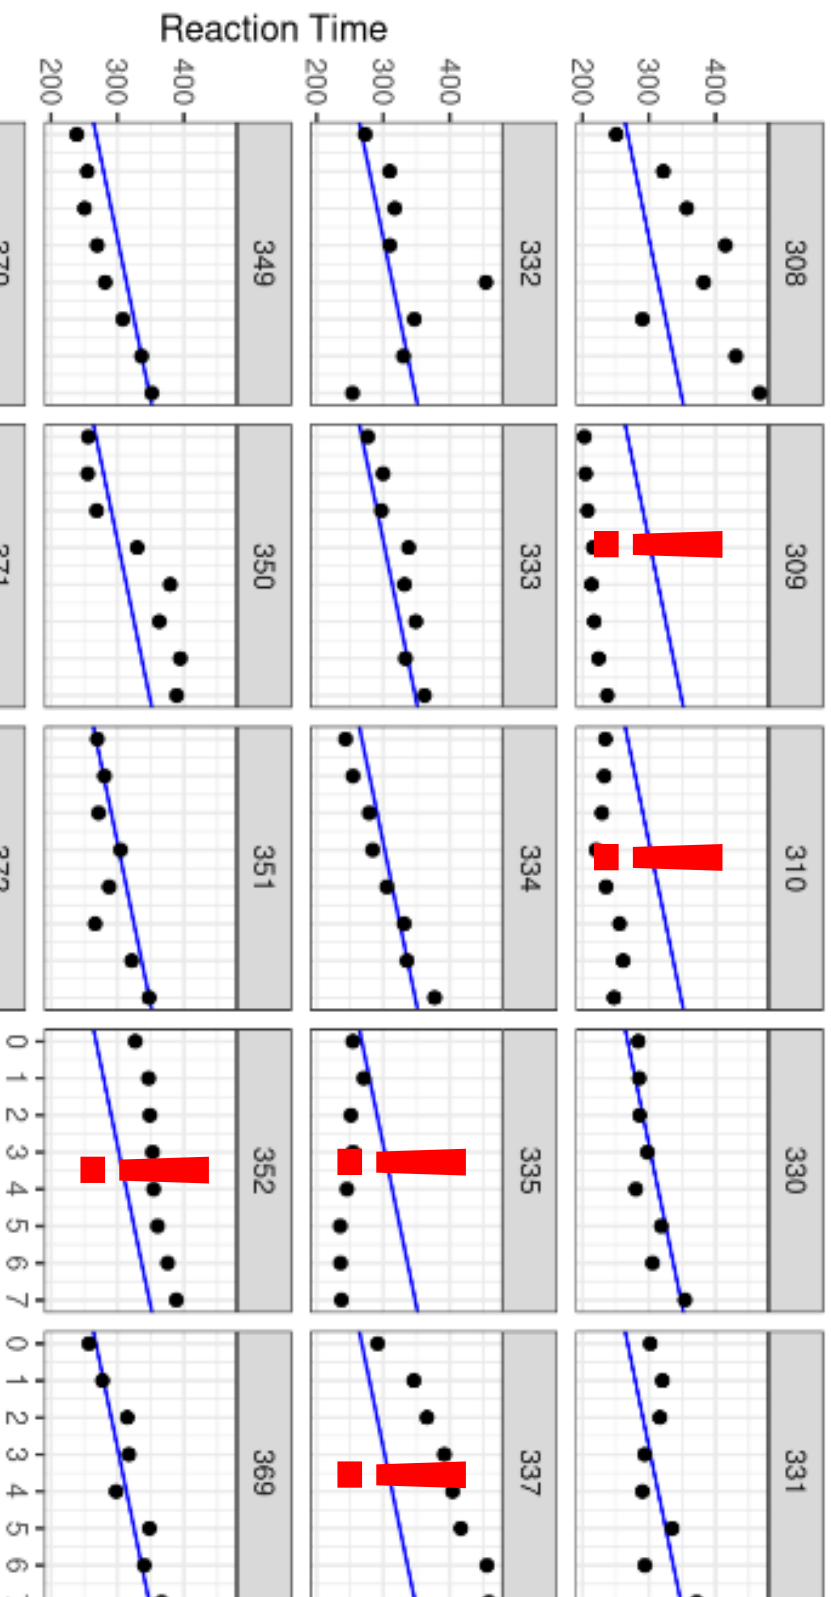 - we model both an average and each level of the categorical predictor (e.g. *subject*)

Figure 5.3: Data from Belenky et al. (2003), showing reaction time at baseline (0) and after each day of sleep deprivation.

(Barr, University of Glasgow)

Reaction Time

Days deprived of sleep (0 = baseline)

COMPLETE POOLING

Call:
lm(formula = Reaction ~ days_deprived, data = sleep2)

Coefficients:
                Estimate
(Intercept)   267.967
days_deprived   11.435

There are some bad fits

(Barr, University of Glasgow)

8

```
lm(formula = Reaction ~ days_deprived + Subject + days_deprived:Subject,
   data = sleep2)
```

```
## Coefficients:
##
##                  Estimate
## (Intercept)      288.2175
## days_deprived     21.6905
## Subject309       -87.9262
## Subject310       -62.2856
## Subject330       -14.9533
## Subject331         9.9658
## Subject332        27.8157
```

... and the remaining 12 subjects

NO POOLING

```
## days_deprived:Subject309  -17.3334
## days_deprived:Subject310  -17.7915
## days_deprived:Subject330  -13.6849
## days_deprived:Subject331  -16.8231
## days_deprived:Subject332  -19.2947
## days_deprived:Subject333  -10.8151
```

... and the remaining 12 subjects

(Barr, University of Glasgow)

Figure 5.5: Data plotted against fits from the no-pooling approach.

NO POOLING

Good fits now:

What are the limits of this model?

(Barr, University of Glasgow)

10

PARTIAL POOLING

Figure 5.6: Data plotted against predictions from a partial pooling approach.

```
Fixed effects:
               Estimate Std. Error t value
(Intercept)     267.967      8.266  32.418
days_deprived    11.435      1.845   6.197
```

`ranef(pp_mod)[["Subject"]]`

	(Intercept)	days_deprived
308	24.4992891	8.6020000
309	-59.3723102	-8.1277534
310	-39.4762764	-7.4292365
330	1.3500428	-2.3845976

```
Linear mixed model fit by REML ['lmerMod']
Formula: Reaction ~ days_deprived + (days_deprived | Subject)
   Data: sleep2
```

(Barr, University of Glasgow)

11

# No pooling vs partial pooling

Figure 5.5: Data plotted against fits from the no-pooling approach.

Figure 5.6: Data plotted against predictions from a partial pooling approach.

Both model the individual variance – but only one is generalisable outside the subject pool
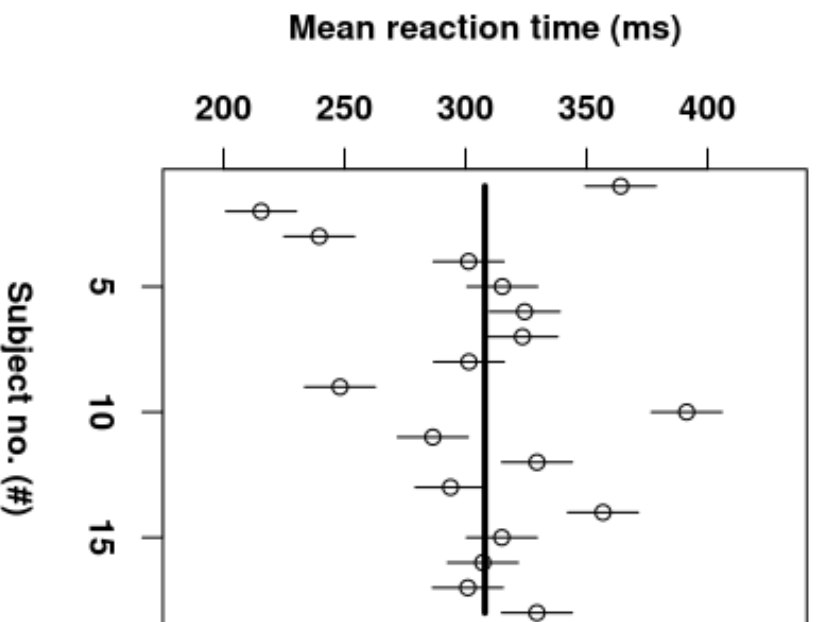
(Barr, University of Glasgow)

12

Partial pooling

(Gelman and Hill, 2006
(12.1))

another scary
looking thing...

$$\hat{\alpha}_j^{multilevel} \approx \frac{\frac{n_j}{\sigma_y^2} \bar{y}_j + \frac{1}{\sigma_\alpha^2} \bar{y}_{all}}{\frac{n_j}{\sigma_y^2} + \frac{1}{\sigma_\alpha^2}}$$
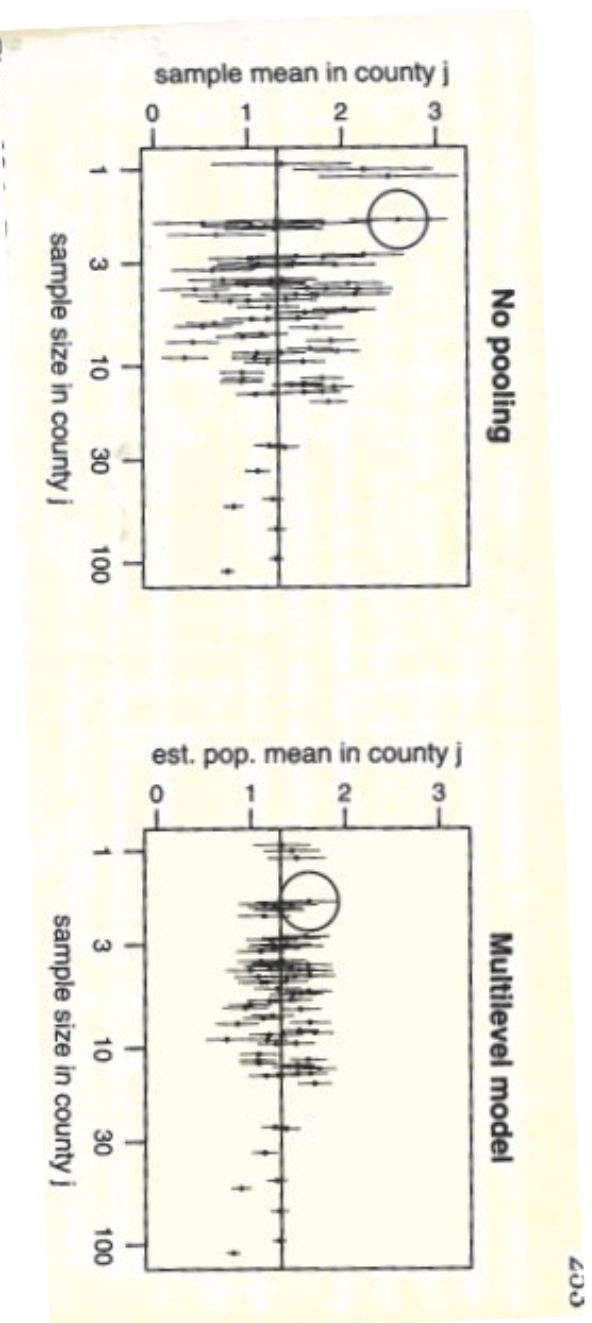
No pooling

Mean reaction time (ms)

Subject no. (#)

Partial pooling

Mean reaction time (ms)

Subject no. (#)

What is the advantage of the partial pooling model?

Now with different sample sizes

(Gelman and Hill, 2006)

What is the advantage of the partial pooling model?

$$\hat{\alpha}^{multilevel} \approx \frac{\dfrac{n_j}{\sigma_y^2}\bar{y}_j + \dfrac{1}{\sigma_\alpha^2}\bar{y}_{all}}{\dfrac{n_j}{\sigma_y^2} + \dfrac{1}{\sigma_\alpha^2}}$$

(Gelman and Hill, 2006)

$\hat{\alpha}_j$: estimated mean for subject j

n_j : sample size for subject j

σ_y^2 : within-subject variance

σ_α^2 : variance around the average

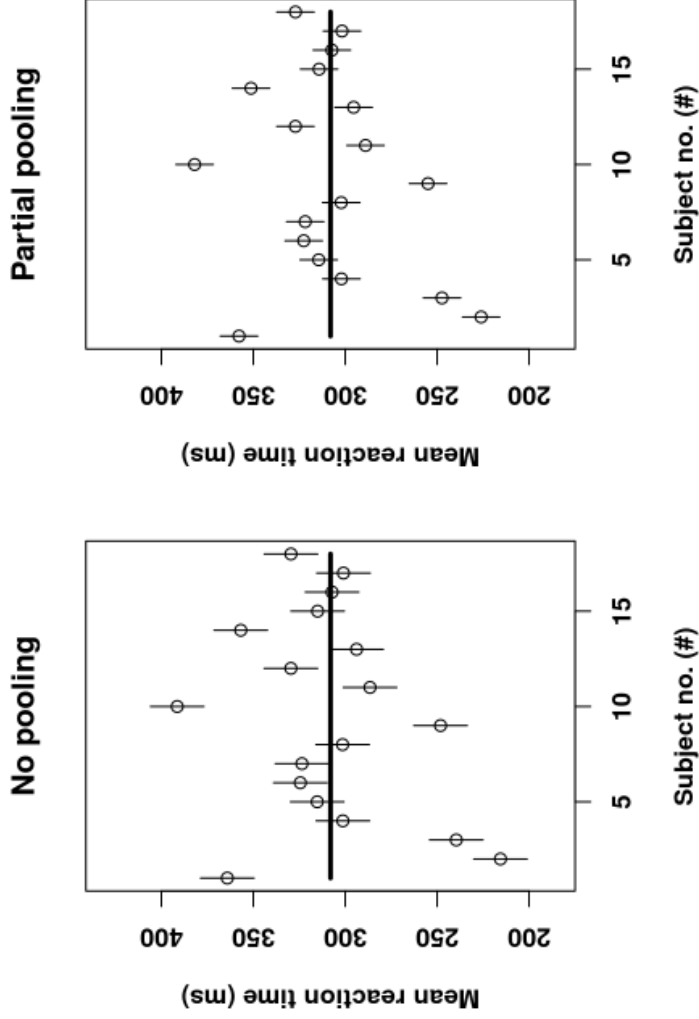$\bar{y}_j$: unpooled estimate of subject j

$\bar{y}_{all}$: the pooled estimate

Discuss in small groups

What happens to the estimated mean, $\hat{\alpha}_j$, when n_j

1) increases?
2) decreases?
3) is 0?
4) goes towards infinity?

Same n for each subject

Different n for each county

(Gelman and Hill, 2006)

```
estimate.multilevel.mean <- function(n.j, sigma.y, sigma.mean, y.j, y.al
l)
{
    alpha <- ((n.j / sigma.y^2) * y.j + (1 / sigma.mean^2) * y.all) /
             ( (n.j / sigma.y^2) + (1 / sigma.mean^2))
    return(alpha)
}
```

"Baseline" plot

```
## "baseline"

sigma.y <- 3
y.j <- 7
sigma.mean <- 1.5
y.all <- 4
ns <- 0:100
```

18

"Baseline" plot

```
## "baseline"

sigma.y <- 3
y.j <- 7
sigma.mean <- 1.5
y.all <- 4
ns <- 0:100
```

Small group effect

```
## small group effect

sigma.y <- 3
y.j <- 7
sigma.mean <- 1.5
y.all <- 1
ns <- 0:100
```

"Baseline" plot

```
## "baseline"
sigma.y <- 3
y.j <- 7
sigma.mean <- 1.5
y.all <- 4
ns <- 0:100
```

Estimated mean

n.samples

Noisy individual effect

```
## noisy individual effect
sigma.y <- 6
y.j <- 7
sigma.mean <- 1.5
y.all <- 4
ns <- 0:100
```

Estimated mean

n.samples

"Baseline" plot

```
sigma.y <- 3
y.j <- 7
sigma.mean <- 1.5
y.all <- 4
ns <- 0:100
```


Noisy group effect

```
## noisy group effect
sigma.y <- 3
y.j <- 7
sigma.mean <- 3.5
y.all <- 4
ns <- 0:100
```

FROM LAST WEEK

Motivation for multilevel modelling:

We want to use all the information in the data while fulfilling the assumptions necessary for the residuals

We can add:

Without letting small or uncertain samples unduly affect our group estimate

$\hat{\alpha}^{multilevel} \approx$

$$\hat{\alpha}_j^{multilevel} \approx \frac{\frac{n_j}{\sigma_y^2} \bar{y}_j + \frac{1}{\sigma_\alpha^2} \bar{y}_{all}}{\frac{n_j}{\sigma_y^2} + \frac{1}{\sigma_\alpha^2}}$$

(Gelman and Hill, 2006 (12.1))

$\hat{\alpha}_j$: estimated mean for subject j

n_j : sample size for subject j

σ_y^2 : within-subject variance

σ_α^2 : variance around the average

$\bar{y}_j$: unpooled estimate of subject j

$\bar{y}_{all}$: the pooled estimate

Revisiting equation 12.1

Why is this not very interesting for purposes of fitting models?

Now with parameters (β)

$$\hat{\alpha}_j \approx \frac{\frac{n_j}{\sigma_y^2}}{\frac{n_j}{\sigma_y^2} + \frac{1}{\sigma_\alpha^2}} \left(\bar{y}_j - \beta \bar{x}_j \right) + \frac{\frac{1}{\sigma_\alpha^2}}{\frac{n_j}{\sigma_y^2} + \frac{1}{\sigma_\alpha^2}} \mu_\alpha$$

$\hat{\alpha}_j$: group level parameters
n_j: sample size for subject j
σ_y^2: within-subject variance
σ_α^2: variance around the average
$\bar{y}_j$: unpooled estimate of subject j
$\bar{y}_{all}$: the pooled estimate
$(\bar{y}_j - \beta \bar{x}_j)$: the unpooled estimate for the subject
μ_α: mean

Now on to generalized linear mixed models ...

Did you learn?

Linear Mixed Effects Models (LMM)

1) Why can it be a good idea to do mixed effects modelling?

2) Understanding the basics of multilevel modelling

   - also known as linear mixed effects modelling

3) Appreciating the difference between the different levels of effects

   – or *random* and *fixed* effects, as they are also called

4) Understanding the concept of pooling (none, complete and partial)

... but let's do a recap of the generalized linear model first

Learning goals
Generalized Linear Mixed Effects Models (GLMM)

1) Understanding that we can extend the scope of our multilevel modelling by using appropriate link functions and data distributions

2) Understanding the multilevel equivalent of the GLM

28

At least four ingredients needed

1) A data vector: $y = (y_1, \ldots, y_n)$

2) Predictors: X and coefficients β, forming a linear predictor $X\beta$

3) A *link function g*: yielding a vector of transformed data $\hat{y} = g^{-1}(X\beta)$ that are used to model the data

4) A data distribution: $p(y \mid \hat{y})$

$$(X\beta = \beta_0 + X_1\beta_1 + \ldots + X_k\beta_k)$$

(Gelman and Hill, 2006,
Chapter 6)

29

Breaking all promises and going back to *mtcars*

1) A data vector: $y = (y_1, ..., y_n)$

```
print(y <- mtcars$am)
```

```
##  [1] 1 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 0 0 0 0 0 1 1 1 1 1 1 1
```

31

2) Predictors: X and coefficients β, forming a linear predictor $X\beta$

```
logistic.model <- glm(am ~ wt + 1, data=mtcars, family='binomial')
X <- model.matrix(logistic.model)
print(head(X))
```

```
##                     (Intercept)    wt
## Mazda RX4                     1 2.620
## Mazda RX4 Wag                 1 2.875
## Datsun 710                    1 2.320
## Hornet 4 Drive                1 3.215
## Hornet Sportabout             1 3.440
## Valiant                       1 3.460
```

```
print(beta.hat <- logistic.model$coefficients)
```

```
## (Intercept)          wt
##    12.04037    -4.02397
```
estimate of our beta hats (intercept and slope)

32

3) A *link function g*: yielding a vector of transformed data $\hat{y} = g^{-1}(X\beta)$ that are used to model the data
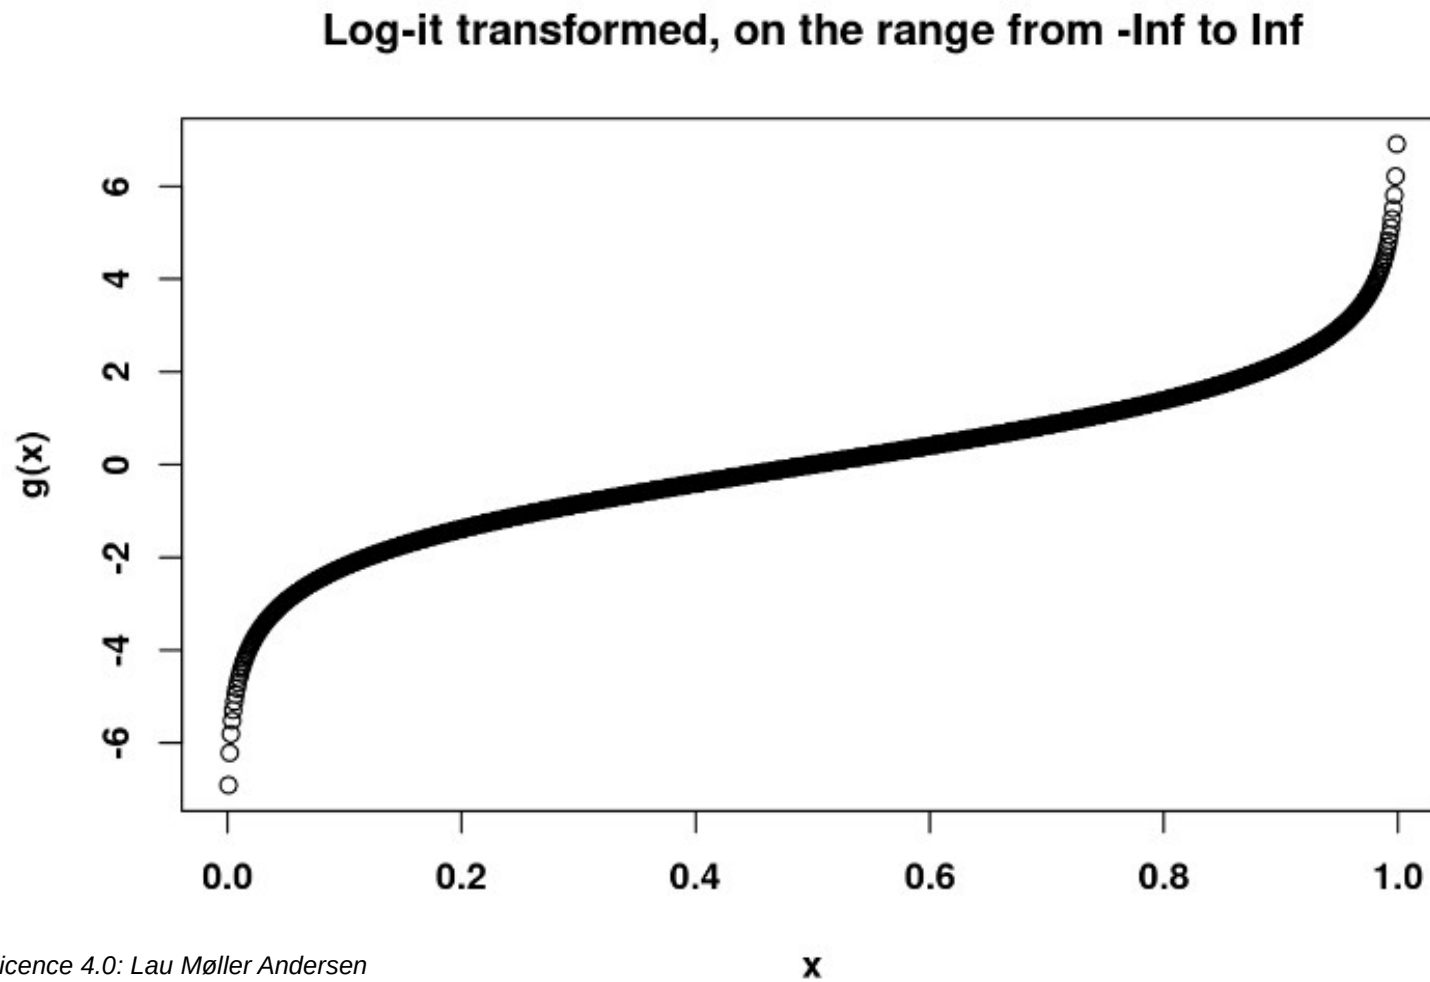
```r
g <- function(x) log(x / (1 - x)) ## logit
inv.g <- function(x) exp(x) / (1 + exp(x)) ##logit-¹
```

probabilities from .001 to 0.999, .001

```
x <- seq(0.001, 0.999, 0.001)
plot(x, main='Original probability data (on the range from 0-1)')
```

Original probability data (on the range from 0-1)

34
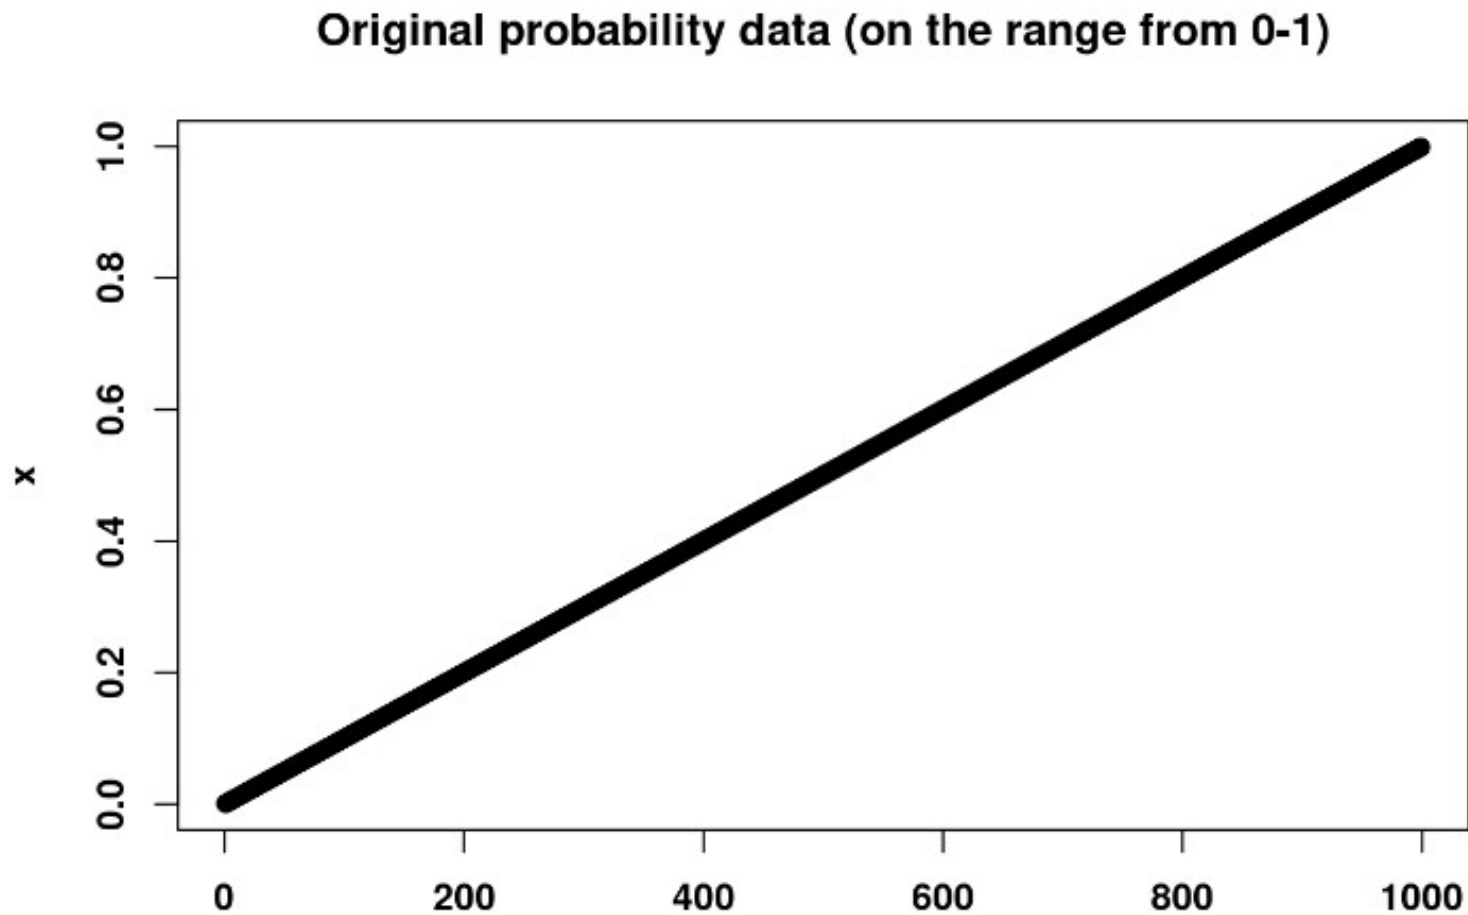
```
plot(x, g(x), main='Log-it transformed, on the range from -Inf to Inf')
```

Log-it transformed, on the range from -Inf to Inf

35

```
plot(x, inv.g(g(x)), main='Back on the original scale')
```

Back on the original scale
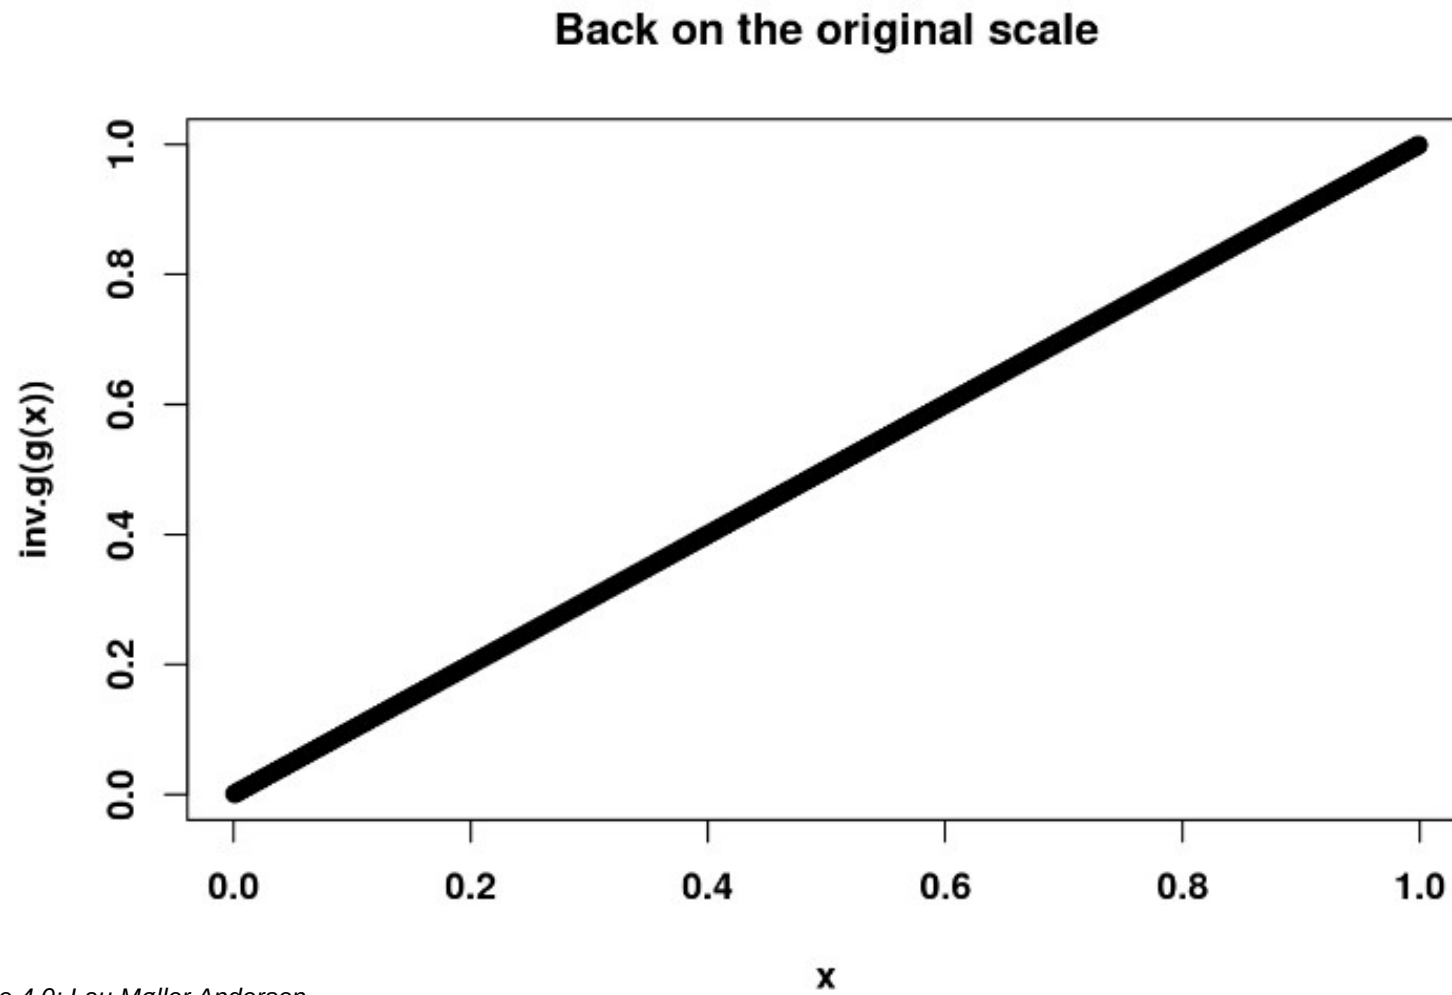
These are the fitted values

```
y.hat <- inv.g(X %*% beta.hat)
print(head(y.hat))
```

```
##                            [,1]
## Mazda RX4             0.8172115
## Mazda RX4 Wag         0.6157283
## Datsun 710            0.9373069
## Hornet 4 Drive        0.2897304
## Hornet Sportabout     0.1415972
## Valiant               0.1320944
```

```
print(head(y.hat - logistic.model$fitted.values))
```

```
##                         [,1]
## Mazda RX4                  0
## Mazda RX4 Wag              0
## Datsun 710                 0
## Hornet 4 Drive             0
## Hornet Sportabout          0
## Valiant                    0
```

These are the linear predictors

linear predictors (no inverse added)

```
y.hat.linear <- X %*% beta.hat

print(head(y.hat.linear - logistic.model$linear.predictors))
```

```
##                    [,1]
## Mazda RX4            0
## Mazda RX4 Wag        0
## Datsun 710           0
## Hornet 4 Drive       0
## Hornet Sportabout    0
## Valiant              0
```

Looks like a "normal" linear regression

Linear predictors

4) A data distribution: $p(y|\hat{y})$

$$\Pr(y=1)=\hat{y}$$

https://en.wikipedia.org/wiki/Binomial_distribution

Some link functions

Usage

```
family(object, ...)

binomial(link = "logit")
gaussian(link = "identity")
Gamma(link = "inverse")
inverse.gaussian(link = "1/mu^2")
poisson(link = "log")
quasi(link = "identity", variance = "constant")
quasibinomial(link = "logit")
quasipoisson(link = "log")
```

Important difference from the general linear model

$$\hat{\beta} = (X^T X)^{-1} X^T Y \textbf{ (least squares solution)}$$

The least squares solution is the maximum likelihood estimation

Important difference from the general linear model

We also make maximum likelihood
estimates for logistic regression,
but there are no analytical solutions
for those

43

Maximum likelihood estimate

Likelihood: $L(\theta \mid O) = \prod_{i=1}^{n} fx(x_i \mid \mu, \sigma^2)$

where μ and σ^2 are parameters describing a normal distribution

θ: the unknown parameters, e.g. $\hat{\beta}$ and $\hat{\sigma}^2$

O: the observations from a given sample

log-likelihood: $l(\theta \mid 0) = \ln(L(\theta \mid O))$

$MLE: \hat{\theta} = \arg \max l(\theta \mid O)$

The general linear mixed model (GLMM)

$$y = X\beta + Zu + \epsilon$$

$y : N\, x\, 1$ column vector
$X : N\, x\, p$ matrix of p predictor variables
$\beta : p\, x\, 1$ column vector of the first level regression coefficients
$Z : N\, x\, q$ design matrix for the q random effects
$u : q\, x\, 1$ column vector of the second-level effects
$\epsilon : N\, x\, 1$ column vector of the residuals

To generalize to non-linear functions
At least four ingredients needed

1) A data vector: $y = (y_1, \dots, y_n)$

2) Predictors: X and coefficients β, forming a linear predictor $X\beta$

3) A *link function* g: yielding a vector of transformed data $\hat{y} = g^{-1}(X\beta)$ that are used to model the data

4) A data distribution: $p(y \mid \hat{y})$

$$(X\beta = \beta_{0j} + X_{1j}\beta_{1j} + \dots + X_{kj}\beta_{kj})$$

This time a $_j$ added to indicate that all of these are modelled at a second level as well

46

Did you learn?

Generalized Linear Mixed Effects Models (GLMM)

1) Understanding that we can extend the scope of our multilevel modelling by using appropriate link functions and data distributions

2) Understanding the multilevel equivalent of the GLM

References

- Gelman, A., Hill, J., 2006. Data Analysis Using Regression and Multilevel/Hierarchical Models. Cambridge University Press.